

08 Shapley Additive Explanations

EP

2026-07-19

SHapley Additive exPlanations (SHAP)

SHAP values were used to interpret two classification models predicting species richness (“high” or “low”): a random forest, using `kernelshap`, and an XGBoost model using TreeSHAP. Both models were configured to explain predicted probability of “high” richness. Because the XGBoost’s categorical predictors (ROFSW and ROFRS) were one-hot encoded during preprocessing, these were aggregated back into their original variables. This was to ensure fair, feature-for-feature comparison between models.

```
library(sf)
library(kernelshap)
library(shapviz)
library(dplyr)
library(tidymodels)
library(stringr)
library(patchwork)
library(targets)
library(ranger)

richness_fit_class <- readRDS("Data/richness_fit_class.rds")
richness_fit_xgb_ND <- readRDS("Data/richness_fit_xgb_ND.rds")
richness_recipe_xgb_ND <- readRDS("Data/richness_recipe_xgb_ND.rds")

predictors <- c("dtm_elev", "avg_imperv", "rofsw_risk", "rofrs_risk",
               "avg_area_ha", "treeht", "weighted_precip", "weighted_temp",
               "dist_to_river", "pct_greenpace")

grid_model_all <- st_read("Data/grid_model.gpkg")

## Reading layer 'grid_cells' from data source
## 'C:\Users\ellsp\OneDrive\Documents\1 Master's Project\RMDs\Data\grid_model.gpkg'
## using driver 'GPKG'
## Simple feature collection with 489 features and 21 fields
## Geometry type: POLYGON
## Dimension: XY
## Bounding box: xmin: 289277.6 ymin: 86966.9 xmax: 461661.2 ymax: 573772.4
## Projected CRS: OSGB36 / British National Grid

df <- grid_model_all %>%
  sf::st_drop_geometry()
```

Random Forest

```
# Extract model & set predictors
rf_model <- richness_fit_class %>%
  extract_fit_engine()
predictors <- df %>%
  dplyr::select(all_of(predictors))

# Set probability of "high"
predictor_function <- function(object, newdata) {
  predict(object, data = newdata)$predictions[, "high"]}

# Calc SHAP values
set.seed(399)

rf_shap <- kernelshap(
  object = rf_model,
  X = predictors,
  pred_fun = predictor_function)
```


|

|

```
# Plot
sv_rf <- shapviz(rf_shap)


p1 <- sv_importance(sv_rf)
p2 <- sv_importance(sv_rf, kind = "beeswarm")

rfcombined <- p1 + p2
#ggsave("rf_shap.png", rfcombined, width = 12, height = 5, dpi = 300)
p1
```



This bar chart displays the SHAP values for the RF model, where the higher the SHAP value the greater the influence that predictor has on predicting “high” richness. **pct_greenpace** is by far the most important predictor in this model, followed by **weighted_temp** and **weighted_precip**. Neither **ROFSW_risk** or **ROFRS_risk** are highly important in this model, but do provide some benefit.

p2



This beeswarm plot not only shows the importance of each predictor but also the direction of its effect and the associated feature value. For `pct_greenpace`, low values are spread evenly across both negative and positive SHAP values, while high values cluster almost entirely on the right with positive SHAP values. This indicates high greenspace strongly increases the predicted probability of high richness. `weighted_temp` shows a similar pattern where low to medium temperatures are associated with probability of low richness, while high temperatures cluster with positive SHAP values. Low precipitation shows the same relationship with low richness. Although ROFSW has a smaller impact on the overall model, there is a clear split where low ROFSW risk clusters at negative SHAP values and high ROFSW with positive SHAP values.

XGBoost

Because of the dummy step in the XGBoost model, a bake step is introduced to ensure the data is processed in the same way the model was. However, `kernelshap` is not needed here since the XGBoost package has TreeSHAP built in, which `shapviz` can use.

```
# Prepare predictor matrix
predictors <- c("dtm_elev", "avg_imperv", "rofsw_risk", "rofrs_risk",
               "avg_area_ha", "treeht", "weighted_precip", "weighted_temp",
               "dist_to_river", "pct_greenpace")

X_baked <- bake(richness_recipe_xgb_ND, new_data = df) %>%
  dplyr::select(-any_of("richness_class"))

# Set up SHAP
xgb_model <- richness_fit_xgb_ND %>%
  extract_fit_engine()
```

```

# Calc TreeSHAP values
sv_xgb <- shapviz(
  object = xgb_model,
  X_pred = as.matrix(X_baked),
  X = X_baked)

sv_xgb$S <- -sv_xgb$S
sv_xgb$baseline <- -sv_xgb$baseline

```

Rejoin ROFSW and ROFRS

Reform the original ROFSW and ROFRS categories for a beeswarm plot comparable with that of the RF.

```

# Map dummy cols to the correct category
shap_cols <- colnames(sv_xgb$S)
parent_vars <- c("rofrs_risk", "rofsw_risk")
col_map <- tibble(col = shap_cols) %>%
  mutate(parent = case_when(
    str_starts(col, "rofrs_risk") ~ "rofrs_risk",
    str_starts(col, "rofsw_risk") ~ "rofsw_risk",
    TRUE ~ col ))

# Sum SHAP values across dummy cols for each observation & aggregate
S <- as.data.frame(sv_xgb$S)

S_agg <- sapply(unique(col_map$parent), function(p) {
  cols <- col_map$col[col_map$parent == p]
  if (length(cols) == 1) {
    S[[cols]]} else {
    rowSums(S[cols])}})

S_agg <- as.data.frame(S_agg)
colnames(S_agg) <- unique(col_map$parent)


# Rebuild
X_agg <- df %>%
  dplyr::select(all_of(names(S_agg)[names(S_agg) %in% c("rofrs_risk", "rofsw_risk")]))
other_cols <- setdiff(colnames(S_agg), c("rofrs_risk", "rofsw_risk"))
X_agg <- cbind(X_baked[other_cols], X_agg)

# Build new shapviz object
sv_xgb_agg <- shapviz(
  object = as.matrix(S_agg),
  X = X_agg)

# Plot
p3 <- sv_importance(sv_xgb_agg)
p4 <- sv_importance(sv_xgb_agg, kind = "beeswarm")

xgbcombined <- p3 + p4
#ggsave("xgboost_shap.png", xgbcombined, width = 12, height = 5, dpi = 300)
p3

```



As seen in the RF model, the top feature is `pct_greenspace`, followed by `weighted_temp` and `weighted_precip`. As before, neither `ROFSW_risk` or `ROFRS_risk` are highly important in this model, but do provide some benefit.

p4




This beeswarm plot shows that the two models largely agree on both the important predictors and the direction of their effects. As in the RF model, `pct_greenspace` shows an even spread of low values but a clear cluster of high values with positive SHAP values. Both `weighted_temp` and `weighted_precip` again show higher values are associated with a higher predicted probability of high richness. ROFSW risk shows the same pattern seen in the RF model, with low risk associated with probability of low richness and vice versa.


Precipitation and Temperature

```
p5 <- sv_dependence(sv_rf, "weighted_precip", color_var = "weighted_temp")
p6 <- sv_dependence(sv_xgb_agg, "weighted_precip", color_var = "weighted_temp")

precip_temp <- p5 + p6
p5
```



p6



```
#ggsave("precip_temp_shap.png", precip_temp, width = 12, height = 5, dpi = 300)
```


Since these two variables are correlated, a dependence plot was generated to examine their combined effect. Low temperature combined with low precipitation is associated with a predicted probability of low richness, whilst high values are associated with a predicted probability of high. This relationship isn't linear and there appears to be a threshold that must be crossed before the prediction shifts from low to high richness.

Elevation vs Distance to River

```
p7 <- sv_dependence(sv_rf, "dtm_elev", color_var = "dist_to_river")
p8 <- sv_dependence(sv_xgb_agg, "dtm_elev", color_var = "dist_to_river")


elevation_river <- p7 + p8
p7
```

```
## Warning: Removed 4 rows containing missing values or values outside the scale range
## ('geom_point()').
```



p8

```
## Warning: Removed 4 rows containing missing values or values outside the scale range
## ('geom_point()').
```



```
#ggsave("elevation_v_river_shap.png", elevation_river, width = 12, height = 5, dpi = 300)
```

These dependence plots shows a more linear relationship between elevation and distance to the nearest river. Proximity to a river is associated with lower elevation and higher SHAP values, indicating a higher predicted probability of high richness at low elevations, near rivers.